

STOCK MARKET TEXT ANALYTICS

USING STOCK MARKET NEWSFEEDS TO PREDICT DAILY GAINS/LOSSES FOR THE DOW JONES INDUSTRIAL AVERAGE

In this project, the goal is to demonstrate the potential of human emotions and feelings through the written word regarding pre-market newsfeeds of the present and recent past. The data used in the project uses a historical dataset of the market open, close, high, low, a target of 1 for gain, and 0 for loss, and lastly daily sentiment analysis scores using the compounded polarity, to help predict the daily market outcome through recent daily and consistent pre-market newsfeeds which begins before the markets open for the day. The idea is that those who publish key topics of what the markets might consider important news, have keen insights and knowledge on the outcome of the markets for the day.

For our perspective, and assumption for this project is to assume that the financial markets represents chaos theory. The concept is that you can beat the market, but you cannot beat the market all the time. The financial markets run on numerous chaotic systems that influences many conditions, which make the market currently (*1) almost impossible to predict. Prediction relies on some form of mathematics in the traditional sense of how the market will perform. However, it is interesting, within Chaos Theory, there is a term which describes human emotion, and in mathematics, this human emotion is call the *strange attraction*. The goal of the project is to measure the accuracy of News Sentiment Scores from Sentiment Analysis in the Data Science world of Text Analytics, to see if this strange attraction, could be in conjunction with more linear and mathematical financial prediction to make something more accurate overall. However this project concludes on looking at this *Strange Attraction*.

The analysis uses a running study of the Daily News Sentiment Index, from a paper and dataset called “Measuring News Sentiment,” by Adam Hale Shapiro, Moritz Sudhof, and Daniel Wilson from the Federal Reserve Bank of San Francisco. (*2). This dataset will be the training dataset that will analyze recent pre-market newsfeeds. The dataset takes Sentiment Scores from several daily news articles, and as of April 2021, it takes from 24 different newspapers, which are the following, Atlanta Journal-Constitution, Boston Herald, Dallas Morning News, Detroit Free Press, Houston Chronicle, New York Post, New York Times, Palm Beach Post, Philadelphia Inquirer, Providence Journal, San Francisco Chronicle, San Jose Mercury News, Seattle Times, St. Louis Post-Dispatch, St. Paul Pioneer Press, Star Ledger, Star Tribune, The Tennessean, Arizona Republic, Arkansas Democrat Gazette, Denver Post, Times-Picayune, Wall Street Journal, and lastly the Washington Post. The entire dataset has a minimum corpus of 238,685 news articles.

First, our test will be taking “Top 5 Things to Know Before the Markets Open” from CNBC’s daily pre-market daily newsfeed, with the following assumptions, CNBC is the top business news network in television, and it is a well-known and main-stream news organization. Even though the Daily

News Sentiment Index uses several newspapers, our assumption is the CNBC's premarket newsfeed is a good general aggregation of the prior, and will assume that it is similar for now.

The idea, we know that Sentiment analysis of Twitter feeds can predict up and down changes in the Dow Jones Industrial Average (DJIA) closing values with an accuracy of 87.6%. However, can we use Premarket Newsfeeds to predict the Dow Jones Industrial Average closing values for the day? If this was possible, and you could get a good accuracy rating with Machine Learning, this will be feasible to make a living, only if it were so easy. This project uses NLTK's Vader Lexicon package that provides Sentiment Analysis Scores for the Premarket Newsfeeds that we are using.

BUSINESS/DATA UNDERSTANDING

There are 3 main datasets used in this project. 1. CNBC '5 things to know before the Stock Market Opens today', newsfeed that was collected over the past 3 weeks and created specifically for this project that was done independently. 2. Stock Market News Headline Sentiment Analysis, by a dataset provided by the Federal Reserve Bank of San Francisco called the "Daily News Sentiment Index," 3. Historical Dow Jones Industrial Average for Open, Close, High, Low, Change, and target of 1 if overall gain, or target 0, if overall loss.

Daily New Sentiment Index, "Measuring News Sentiment" *Federal Reserve Bank of San Francisco'

The Daily News Sentiment Index, which is from "Measuring News Sentiment" from the article from the Federal Reserve Bank of San Francisco, using an average since 1980 of several newspapers and judges around 25 of them consistently through this period of time.

	$s_a^i = f_{t(a)}^i + f_{p(a),j(a)}^i + \varepsilon_a^i$	
Where		
o is the positivity score for article	s_a^i	
o is a sample-day (t) fixed effect	$f_{t(a)}^i$	
o is a newspaper*type fixed effect	$f_{p(a),j(a)}^i$	

What this model does is take lexicons and use its own model that looks at the sentiment in economic news, not just daily newsfeeds during pre-market or intraday market. It uses the Loughran and McDonald dictionary and Harvard General Inquirer dictionary. For NLTK, VADER package, we will use the default, however it is understood that these two are different. This model does score with weighted averages that we do not use in VADER for the TEST dataset. It would be interesting to see if we can improve the accuracy of the prediction scores and values if we used a weight on the VADER score that we later use in the TEST dataset. However, please note that this project does not yet do that.

Below is one of the dictionaries and lexicon scores and rules for the Sentiment Analysis Scoring used for the Daily News Sentiment Index.

Word	Sequence	Word Count	Word Proportion	Average Proportion	Std Dev	Doc Count	Negative	Positive	Uncertainty	Litigious	Constrained	Superfluous	Interesting	Modal	Irr_Verb	Harvard_IV	Syllables	Source
AARDVARK	1	275	1.60E-08	1.31E-08	3.67E-06	82	0	0	0	0	0	0	0	0	0	0	0	2 12of12inf
AARDVARKS	2	3	1.75E-10	1.03E-11	1.01E-08	1	0	0	0	0	0	0	0	0	0	0	0	2 12of12inf
ABACI	3	8	4.66E-10	1.47E-10	6.40E-08	7	0	0	0	0	0	0	0	0	0	0	0	3 12of12inf
ABACK	4	6	3.50E-10	1.76E-10	7.21E-08	6	0	0	0	0	0	0	0	0	0	0	0	2 12of12inf
ABACUS	5	6,729	3.92E-07	3.75E-07	3.45E-05	845	0	0	0	0	0	0	0	0	0	0	0	3 12of12inf
ABACUSES	6	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	4 12of12inf

Example of the DNSI Code below with the Loughran and McDonald Dictionary:

```
def load_lm_lexicon():
    # Loughran McDonald
    fn = os.path.join(lexicon_dir, "LoughranMcDonald_2016.csv")
    reader = csv.DictReader(open(fn))
    words2weights = {}
    for r in reader:
        pos_score = 1. if r['Positive'] != "0" else 0.
        neg_score = 1. if r['Negative'] != "0" else 0.
        sentiment_score = pos_score - neg_score
        w = r['\uffffWord'].lower() # weird header in this file
        #w = r['Word'].lower()
        if sentiment_score:
            words2weights[w] = sentiment_score
    return words2weights
```

Vader Lexicon using NLTK in Python with CNBC ‘5 things to know before the Stock Market Opens today’,

VADER Formulation for the Compounded Polarity Score.

I collected 15 days’ worth of the Daily Premarket Newsfeed. It would be much more interesting if more were collected, however the automated idea of web scrapping, which I provide as interesting code that I did and tried with this project is provided. However, the result was not conclusive, or simply did not work with the type of web pages were are trying to get data from.

$$\frac{x}{\sqrt{x^2 + \alpha}}$$

DATA PREPERATION

Data preparation phase with our good old friend CRISP, I went through the motions of tokenization, look at a corpus vs. document, looking at bigrams, a bag of words, reviewing the term frequency, removing stopwords, punctuation, numbers, lemmatization, conversion of text to lower case. One area that was planned, and was done, but was not necessary, after completion, and was a surprising result was that VADER and NLTK does a nice job of preprocessing the text, without doing the necessary steps one-by-one.

Merging of the 2 datasets and preparing the project dataset as the TEST dataset. A massive V-LOOKUP in Excel was used to merge the two historical datasets which will be used as the training dataset for the machine learning model that we will use for the prediction. However, we are using the text Sentiment Scores provided from the 'News Sentiment Dataset,' which is the Daily News Sentiment Index from the Federal Reserve Bank of San Francisco. Removal of Dates on Stock Market Holidays from the News Sentiment Dataset (Daily News Sentiment Index). The dataset came with everyday of the week since 1980. This was trimmed and merged other datasets to make it complete. Data was removed for weekends and the following years with the following Holidays, 2012 - 2023 Stock Market Holidays

Data preparation phase with our good old friend CRISP, I went through the motions of tokenization, look at a corpus vs. document, looking at bigrams, a bag of words, reviewing the term frequency, removing stopwords, punctuation, numbers, lemmatization, conversion of text to lower case. Surprisingly, which was overlooked, was that NLTK and the Vader package and its Sentiment Polarity Score automatically does a nice job of preprocessing the data, which makes preprocessing not necessary with NLTK and Vader. I also added Open Close, High Low, Change, and a target to the DNSI. Also created my own dataset with CNBCs premarket newsfeeds.

News Sentiment Dataset, created => 07/01/2012 – 02/24/2023 Dataset

This was created by Merging the Daily News Sentiment Index, which only had a date and Sentiment Polarity Score. Used Historical Datasets from MarketWatch, which only lets you pull 1 year's worth of Dow Jones Industrial Average data.

date	News Sentiment	Open	Close	High	Low	Change	target
07/02/12	-0.219829142	12879.71	12871.39	12902.12	12795.48	-8.32	0
07/03/12	-0.228132448	12868.06	12943.82	12946.2	12845.28	75.76	1
07/05/12	-0.207128756	12941.85	12896.67	12961.3	12852.24	-45.18	0
07/06/12	-0.208013396	12889.4	12772.47	12889.4	12702.99	-116.93	0
07/09/12	-0.195998693	12772.02	12736.29	12772.02	12686.57	-35.73	0
07/10/12	-0.189461006	12733.87	12653.12	12830.29	12606.91	-80.75	0
07/11/12	-0.195956946	12653.04	12604.53	12661.97	12534.33	-48.51	0
07/12/12	-0.214600266	12602.71	12573.27	12630.64	12492.25	-29.44	0
07/13/12	-0.234180128	12573.73	12777.09	12784.73	12573.04	203.36	1
07/16/12	-0.232607908	12776.33	12727.21	12779.58	12690.05	-49.12	0

This was accomplished using a massive VLOOKUP with another dataset. Example of how a massive VLOOKUP can be used in excel below if it might be useful to anyone else.

1	2	3	4	5	6	7	8	9	10
date	News Sentiment	target	Top1	Top2	Top3	Top4	Top5	Top6	Top7
1/1/2000	0.281809731	#N/A	#N/A	#N/A	#N/A	#N/A	#N/A	#N/A	#N/A
1/2/2000	0.29749183	#N/A	#N/A	#N/A	#N/A	#N/A	#N/A	#N/A	#N/A
1/3/2000	0.308622721		0 A 'hindrance to	Scorecard	Hughes' in	Jack gets h	Chaos as I	Depleted L	Hungry Sp
1/4/2000	0.291098159		0 Scorecard	The best la	Leader: Ge	Cheerio, b	The main r	Has Cubie	Has Cubie
1/5/2000	0.305734373		0 Coventry caugl	United's ri	Thatcher is	Police hel	Tale of Tr	England or	Pakistan re
1/6/2000	0.318497865		1 Pilgrim knows l	Thatcher f	McIlroy ca	Leicester k	United bra	Auntie bac	Shoaib app
1/7/2000	0.309894685		1 Hitches and Hc	Beckham c	Breast can	Alan Parke	Guardian r	Hollywooc	Ashes and
1/8/2000	0.301847351	#N/A	#N/A	#N/A	#N/A	#N/A	#N/A	#N/A	#N/A
1/9/2000	0.286981228	#N/A	#N/A	#N/A	#N/A	#N/A	#N/A	#N/A	#N/A
1/10/2000	0.280951615		1 Fifth round dra	BBC unveil	Second Div	European	Third Divis	Welfare cc	Ferguson p
1/11/2000	0.278537152		1 Man Utd 2 - 0	How North	Buoyant B	Tranmere	United sit	Queen's Pa	Waugh hit
1/12/2000	0.291446413		0 Newcastle see	Liverpool	Highlander	Edwards' p	Chelsea ga	Taylor sett	Tenth top
1/13/2000	0.302617636		1 Bungling offic	And in the	United put	England ag	Donald po	Adams sta	Money mc
1/14/2000	0.304831842		1 Pompey plump	Roma und	Prenton Pa	OK, I didn't	Chelsea tu	Top storey	West Indie
1/15/2000	0.341403395	#N/A	#N/A	#N/A	#N/A	#N/A	#N/A	#N/A	#N/A
1/16/2000	0.329100193	#N/A	#N/A	#N/A	#N/A	#N/A	#N/A	#N/A	#N/A
1/17/2000	0.331310601	#N/A	#N/A	#N/A	#N/A	#N/A	#N/A	#N/A	#N/A

1	2	3	4	5	6	7	8	9
date	News Sentiment	target	Top1	Top2	Top3	Top4	Top5	Top6
36526	0.281809730871954	=VLOOKUP(\$A3,Text!\$A:\$AA,Text!B\$1,0)	=VLOOKUP(\$A3,Text!\$A:\$AA,Text!B\$1,0)	=VLOOKUP(\$A3,Text!\$A:\$AA,Text!B\$1,0)	=VLOOKUP(\$A3,Text!\$A:\$AA,Text!B\$1,0)	=VLOOKUP(\$A3,Text!\$A:\$AA,Text!B\$1,0)	=VLOOKUP(\$A3,Text!\$A:\$AA,Text!B\$1,0)	=VLOOKUP(\$A3,Text!\$A:\$AA,Text!B\$1,0)
36527	0.297491829811852	=VLOOKUP(\$A4,Text!\$A:\$AA,Text!B\$1,0)	=VLOOKUP(\$A4,Text!\$A:\$AA,Text!B\$1,0)	=VLOOKUP(\$A4,Text!\$A:\$AA,Text!B\$1,0)	=VLOOKUP(\$A4,Text!\$A:\$AA,Text!B\$1,0)	=VLOOKUP(\$A4,Text!\$A:\$AA,Text!B\$1,0)	=VLOOKUP(\$A4,Text!\$A:\$AA,Text!B\$1,0)	=VLOOKUP(\$A4,Text!\$A:\$AA,Text!B\$1,0)
36528	0.308622721159146	=VLOOKUP(\$A5,Text!\$A:\$AA,Text!B\$1,0)	=VLOOKUP(\$A5,Text!\$A:\$AA,Text!B\$1,0)	=VLOOKUP(\$A5,Text!\$A:\$AA,Text!B\$1,0)	=VLOOKUP(\$A5,Text!\$A:\$AA,Text!B\$1,0)	=VLOOKUP(\$A5,Text!\$A:\$AA,Text!B\$1,0)	=VLOOKUP(\$A5,Text!\$A:\$AA,Text!B\$1,0)	=VLOOKUP(\$A5,Text!\$A:\$AA,Text!B\$1,0)
36529	0.291098158985108	=VLOOKUP(\$A6,Text!\$A:\$AA,Text!B\$1,0)	=VLOOKUP(\$A6,Text!\$A:\$AA,Text!B\$1,0)	=VLOOKUP(\$A6,Text!\$A:\$AA,Text!B\$1,0)	=VLOOKUP(\$A6,Text!\$A:\$AA,Text!B\$1,0)	=VLOOKUP(\$A6,Text!\$A:\$AA,Text!B\$1,0)	=VLOOKUP(\$A6,Text!\$A:\$AA,Text!B\$1,0)	=VLOOKUP(\$A6,Text!\$A:\$AA,Text!B\$1,0)
36530	0.305734373198226	=VLOOKUP(\$A7,Text!\$A:\$AA,Text!B\$1,0)	=VLOOKUP(\$A7,Text!\$A:\$AA,Text!B\$1,0)	=VLOOKUP(\$A7,Text!\$A:\$AA,Text!B\$1,0)	=VLOOKUP(\$A7,Text!\$A:\$AA,Text!B\$1,0)	=VLOOKUP(\$A7,Text!\$A:\$AA,Text!B\$1,0)	=VLOOKUP(\$A7,Text!\$A:\$AA,Text!B\$1,0)	=VLOOKUP(\$A7,Text!\$A:\$AA,Text!B\$1,0)

3	4
target	Top1
=VLOOKUP(\$A3,Text!\$A:\$AA,Text!B\$1,0)	=VLOOKUP(\$A3,Text!\$A:\$AA,Text!C\$1,0)
=VLOOKUP(\$A4,Text!\$A:\$AA,Text!B\$1,0)	=VLOOKUP(\$A4,Text!\$A:\$AA,Text!C\$1,0)
=VLOOKUP(\$A5,Text!\$A:\$AA,Text!B\$1,0)	=VLOOKUP(\$A5,Text!\$A:\$AA,Text!C\$1,0)
=VLOOKUP(\$A6,Text!\$A:\$AA,Text!B\$1,0)	=VLOOKUP(\$A6,Text!\$A:\$AA,Text!C\$1,0)
=VLOOKUP(\$A7,Text!\$A:\$AA,Text!B\$1,0)	=VLOOKUP(\$A7,Text!\$A:\$AA,Text!C\$1,0)

1	2	3	4	5	6	7	8	9	10	11	12	13	14
Date	Label	Top1	Top2	Top3	Top4	Top5	Top6	Top7	Top8	Top9	Top10	Top11	Top12
1/3/2000	0	A 'hindran	Scorecard	Hughes' in	Jack gets h	Chaos as N	Depleted L	Hungry Sp	Gunners sc	Derby rais	Southgate	Hammers	Saints part
1/4/2000	0	Scorecard	The best la	Leader: Ge	Cheerio, b	The main r	Has Cubie	Has Cubie	Has Cubie	Hopkins 'f	Has Cubie	A tale of t	I say what
1/5/2000	0	Coventry c	United's ri	Thatcher is	Police hel	Tale of Tr	England or	Pakistan re	Cullinan cc	McGrath p	Blair Witcl	Pele turns	Party divid
1/6/2000	1	Pilgrim kn	Thatcher f	McIlroy ca	Leicester k	United bra	Auntie bac	Shoaib app	Hussain hu	England's c	Revenge is	Our choice	Profile of f
1/7/2000	1	Hitches an	Beckham c	Breast can	Alan Parke	Guardian r	Hollywooc	Ashes and	Whingers	Alan Parke	Thuggery,	Met faces	Everton fa

CNBC's 5 Things to Know Before the Markets Open Today, Created Dataset

Below is how I prepared this dataset, which will be the TEST dataset for Machine Learning. It would be a much more manual ML process by measuring the real accuracy of the TEST, rather than do an 80/20 split. This seems much more interesting comparing the results side by side from the predicitons in Machine Learning versus the results in reality.


```
[ ] 1 print("Sentiment scores:")
2 print(scores)

Sentiment scores:
('neg': 0.003, 'neu': 0.842, 'pos': 0.07, 'compound': -0.5494)

[ ] 1 ## 03/16/2023
2 ## url = 'https://www.cnbc.com/2023/03/16/5-things-to-know-before-the-stock-market-opens-thursday-march-16.html'
3
4 text14 = "1. What's next in this wild week? Two more trading days to go in this chaotic week that's been driven by turmoil in banks. Credit Suisse, the current poster child for the upheaval, took up the Swiss central bank's offer of a lifeline. Will that calm

[ ] 1 text14 = text14.lower()

[ ] 1 scores = sia.polarity_scores(text14)

[ ] 1 print("Sentiment scores:")
2 print(scores)

Sentiment scores:
('neg': 0.103, 'neu': 0.782, 'pos': 0.115, 'compound': 0.5159)

[ ] 1 ## 03/17/2023
2 ## url = 'https://www.cnbc.com/2023/03/17/5-things-to-know-before-the-stock-market-opens-friday-march-17.html'
3
4 text15 = "1. Coming out ahead it looks like major stock averages could come out ahead this week, on pace for a winning week despite the turmoil in the global banking sector. Through Thursday, the Dow Jones Industrial Average is up 1.06%, the S&P 500 has risen

[ ] 1 text15 = text15.lower()

[ ] 1 scores = sia.polarity_scores(text15)

[ ] 1 print("Sentiment scores:")
2 print(scores)

Sentiment scores:
('neg': 0.045, 'neu': 0.892, 'pos': 0.064, 'compound': 0.7341)
```

I then took the Sentiment Analysis scores, which is the compounded score, and added it to the dataset as the results/scores for the TEST dataset, which we will be looking at its accuracy from reality.

MODELING

The Modeling was done with the intention of a target variable, a binary result, which could easily be used for Logistic Regression, or more advanced Classification, however our PyCaret looks at the best models for us.

TRAINING –

The Training Dataset was primarily the Daily News Sentiment Index Score, including the Date, Open, Close, High, Low, Change, and target variable. The target variable is 1 for if the closing value was positive and the target variable 0 is if the closing value was negative for the Dow Jones Industrial Average, respectfully.

TEST

The TEST dataset was the dataset that I put together on my own from the VADER Sentiment Polarity Compounded Scores, alongside the actual/reality of the stock market results for those days. We use the TEST dataset to see how accurately it can predict reality on the daily CNBC Premarket Newsfeed.

```
[ ] 1 import pandas as pd

[ ] 1 from google.colab import files
2 uploaded = files.upload()

[ ] 1 from google.colab import files
2 uploaded = files.upload()

[ ] 1 train = pd.read_csv('train_text.csv', usecols=['News Sentiment', 'Open', 'Close', 'High', 'Low', 'Change', 'target'])
2 test = pd.read_csv('test_text.csv', usecols=['News Sentiment', 'Open', 'Close', 'High', 'Low', 'Change'])

[ ] 1 from pycaret.classification import *
2 s = setup(data=train, target='target')
```

	Description	Value
0	Session id	8390
1	Target	target
2	Target type	Binary
3	Original data shape	(2679, 7)
4	Transformed data shape	(2679, 7)
5	Transformed train set shape	(1875, 7)
6	Transformed test set shape	(804, 7)
7	Numeric features	6
8	Preprocess	<input checked="" type="checkbox"/>
9	Imputation type	simple
10	Numeric imputation	mean
11	Categorical imputation	mode
12	Fold Generator	StratifiedKFold
13	Fold Number	10
14	CPU Jobs	-1
15	Use GPU	False
16	Log Experiment	False

We include all the necessary columns for Machine Learning, however we stayed away from the dates. Not that dates would cause time series, this is not a time series project.

EVALUATION

```
best = compare_models()
```

	Model	Accuracy	AUC	Recall	Prec.	F1	Kappa	MCC	TT (Sec)
dt	Decision Tree Classifier	1.00000	1.00000	1.00000	1.00000	1.00000	1.00000	1.00000	
rf	Random Forest Classifier	1.00000	1.00000	1.00000	1.00000	1.00000	1.00000	1.00000	
ada	Ada Boost Classifier	1.00000	1.00000	1.00000	1.00000	1.00000	1.00000	1.00000	
gbc	Gradient Boosting Classifier	1.00000	1.00000	1.00000	1.00000	1.00000	1.00000	1.00000	
xgboost	Extreme Gradient Boosting	1.00000	1.00000	1.00000	1.00000	1.00000	1.00000	1.00000	
lightgbm	Light Gradient Boosting Machine	1.00000	1.00000	1.00000	1.00000	1.00000	1.00000	1.00000	
lr	Logistic Regression	0.9984	1.00000	0.9990	0.9980	0.9985	0.9968	0.9968	
et	Extra Trees Classifier	0.9691	0.9959	0.9722	0.9704	0.9712	0.9378	0.9380	
lda	Linear Discriminant Analysis	0.9088	0.9943	0.9940	0.8599	0.9216	0.8141	0.8276	
ridge	Ridge Classifier	0.9077	0.0000	0.9930	0.8590	0.9207	0.8119	0.8254	
nb	Naive Bayes	0.8730	0.9905	0.9921	0.8140	0.8938	0.7401	0.7638	
knn	K Neighbors Classifier	0.8688	0.9401	0.8955	0.8652	0.8797	0.7354	0.7367	
qda	Quadratic Discriminant Analysis	0.8565	0.9806	0.9761	0.8131	0.8826	0.7047	0.7360	
svm	SVM - Linear Kernel	0.7806	0.0000	0.8306	0.8546	0.7717	0.5545	0.6095	
dummy	Dummy Classifier	0.5360	0.5000	1.00000	0.5360	0.6979	0.0000	0.0000	

It was very interesting as how the Decision Tree Classifier, which is one of more simple Data Model Algorithms is selected as one of the best. From previous work, the more complex and computational intense algorithms such as Gradient Boosting seem to always be one of the best, alongside Random Forest.

Once the results were looked at and analyzed, the current model and how the data was prepared only provided a .5333 Accurate Result, even though Open, Close, High, Low, Change, and Target was used in the training. Did this modeling make sense to begin with? The simple answer is no. This is only slightly better than a 50/50 toss-up. However this project is not about Bayesian Inference.

Upon Reflection of the results, there were a couple issues of the modeling that did not make sense. 1. If I have Open, Close, High, Low, and Change for the day, how could that make sense, as we already have the data that tells us what happened, how would those values be useful in the training if we are trying to predict? Also, the Daily News Sentiment Index Score is not the same as the default VADER Sentiment Compounded Score. Changes must be made. The first line is the results, the second line is actually what happened, therefore we were 8/15, only 53% accurate.

```
verbose=False), 'train_text.csv.pkl')

[ ] 1 print(predictions)

0 32906.168156 32889.089844 33189.281258 32814.179688 -17.070000
1 32873.468750 32656.699219 32873.468750 32636.429688 -216.770004
2 32656.369141 32661.839844 32746.150391 32500.710938 5.470000
3 32780.968750 33083.570312 33083.449219 32665.849609 222.600006
4 33076.328125 33390.968750 33485.820312 33088.418156 314.640015
5 33425.320312 33431.441406 33572.210750 33383.468750 6.120000
6 33428.368594 32856.460938 33453.250000 32838.216938 -571.849976
7 32872.078125 32798.398438 32983.441406 32612.699219 -73.680000
8 32876.828125 32254.859375 32990.460938 32190.599609 -621.969971
9 32105.148625 31989.640625 32422.099609 31783.418156 -275.560000
10 31819.029688 31819.140625 32240.349609 31624.802141 -8.790000
11 32855.269062 32152.400391 32306.589844 31895.400391 100.110001
12 31759.869141 31874.570312 31906.470703 31429.820312 114.699987
13 31827.650391 32246.550781 32281.609375 31571.460938 418.809984
14 32217.320312 31861.980469 32217.320312 31728.699219 -355.339996

News Sentiment prediction_label prediction_score
0 -0.6166 0 1.0
1 -0.2023 0 1.0
2 0.7061 1 1.0
3 0.5859 1 1.0
4 0.3612 1 1.0
5 -0.9584 1 1.0
6 -0.9535 0 1.0
7 -0.6788 0 1.0
8 -0.9830 0 1.0
9 -0.4939 0 1.0
10 0.9990 0 1.0
11 -0.9601 1 1.0
12 -0.8494 1 1.0
13 0.5159 1 1.0
14 0.7341 0 1.0

[ ] 1 predictions.describe()

[ ] 1 predictions

1 HHH 0,1,1,0,0,0,0,1,1,1,0,0
2 HHH 0,0,1,1,1,0,0,0,0,1,1,0
3 HHH 1,0,1,1,0,0,1,1,0,0,1,0,0
4
5 8/15

0.5333333333333333

This is the conclusion of the Model. The conclusion is simple, that using CNBC's Pre-Market Newsfeed, '5 Things to Know Before the Markets
Open' since 02/27/2023, which has been a very negative news cycle, which has been very bad news about the economy potentially. Therefore
```

MODELING 2

We want down to a more basic model with the training and test datasets. I removed the Open, Close, High, Low, Kept daily change, and target. For the test dataset I removed Open, Close, High, Low, Kept daily change, but have the daily change precede to the previous day. The logic is that the premarket newsfeeds feeds off the previous days emotions, and results.

TRAIN

date	News Sentiment	Change	target
02/24/23	-0.033027543	-182.27	0
02/23/23	-0.031600946	-21.48	0
02/22/23	-0.020291626	-124.24	0
02/21/23	-0.014879815	-570.1	0
02/17/23	0.00248806	149.68	1
02/16/23	0.024803413	-295.24	0
02/15/23	0.021685863	119.42	1
02/14/23	0.018055353	-104.82	0
02/13/23	-0.005527485	358.54	1
02/10/23	-0.040999878	197.73	1

TEST

Notice how I shifted the Daily Change up 1 cell to reflect the previous day in the TEST. This logic is used as an assumption of reality that the daily Premarket Newsfeeds, feeds partially from the results of the previous day. In the end, the prediction values were much more accurate.

Date	News Sentiment	Change
02/27/23	-0.6166	-182.27
02/28/23	0.2023	-17.07
03/01/23	0.7061	-216.77
03/02/23	0.5859	5.47
03/03/23	0.3612	222.6
03/06/23	-0.9584	314.64
03/07/23	-0.9535	6.12
03/08/23	-0.6788	-571.85
03/09/23	-0.983	-73.68
03/10/23	-0.4939	-621.97
03/13/23	0.999	-275.5
03/14/23	-0.9601	-0.79
03/15/23	-0.8494	100.11
03/16/23	0.5159	114.7
03/17/23	0.7341	418.9

```
[ ] 1 import pandas as pd
```

```
[ ] 1 from google.colab import files
2 uploaded = files.upload()
```

No file chosen Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.
Saving train_text_v2.csv to train_text_v2.csv

```
[ ] 1 from google.colab import files
2 uploaded = files.upload()
```

No file chosen Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.
Saving test_text_v2.csv to test_text_v2.csv

```
[ ] 1 train = pd.read_csv('train_text_v2.csv', usecols=['News Sentiment', 'Change', 'target'])
2 test = pd.read_csv('test_text_v2.csv', usecols=['News Sentiment', 'Change'])
```

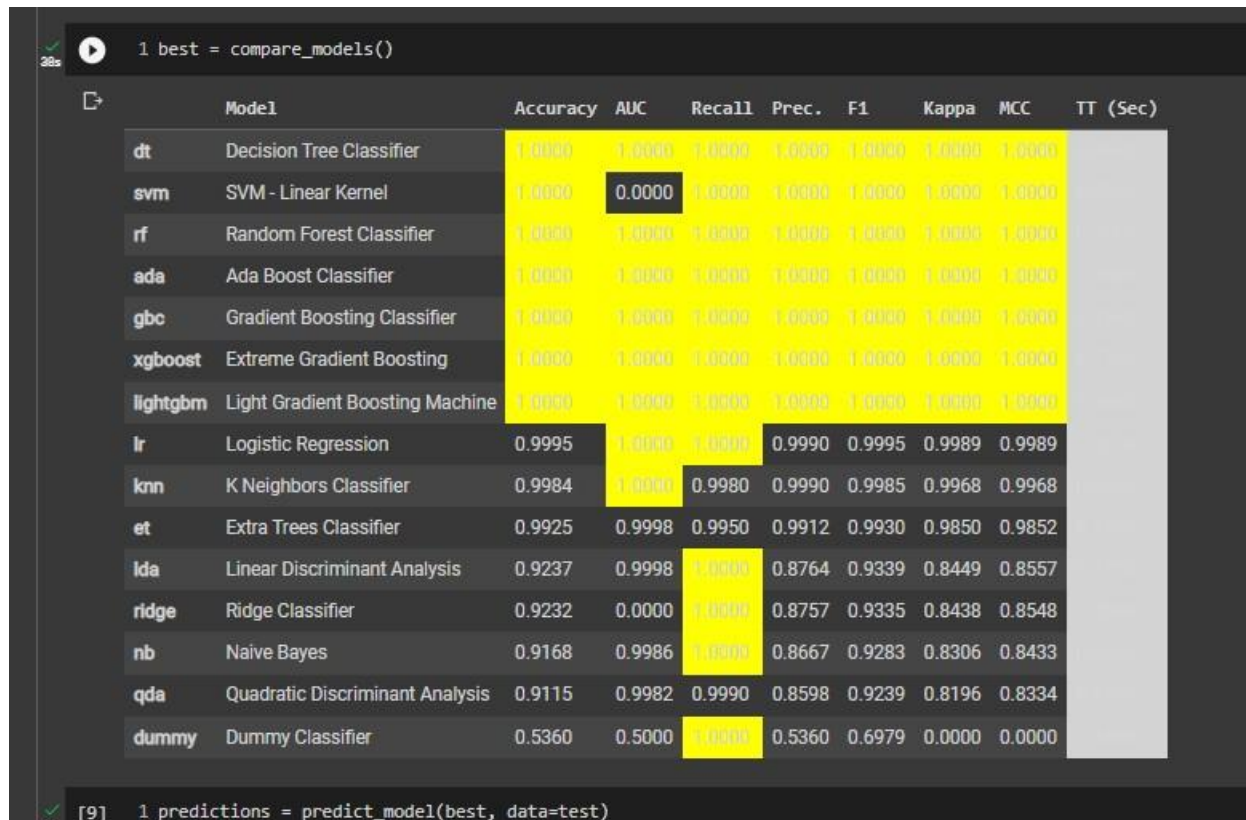
```
1 from pycaret.classification import *
2 s = setup(data=train, target='target')
```

	Description	Value
0	Session id	8644
1	Target	target
2	Target type	Binary
3	Original data shape	(2679, 3)
4	Transformed data shape	(2679, 3)
5	Transformed train set shape	(1875, 3)
6	Transformed test set shape	(804, 3)
7	Numeric features	2
8	Preprocess	yes
9	Imputation type	simple
10	Numeric imputation	mean
11	Categorical imputation	mode
12	Fold Generator	StratifiedKFold
13	Fold Number	10
14	CPU Jobs	-1
15	Use GPU	False
16	Log Experiment	False
17	Experiment Name	clf-default-name
18	USI	fd06

```
[ ] 1 best = compare_models()
```

EVALUATION 2

The second evaluation is has provided much more interesting and successful results, which you would expect from the first evaluation, however many lessons were learned.



	Model	Accuracy	AUC	Recall	Prec.	F1	Kappa	MCC	TT (Sec)
dt	Decision Tree Classifier	1.0000	1.0000	1.0000	1.0000	1.0000	1.0000	1.0000	
svm	SVM - Linear Kernel	1.0000	0.0000	1.0000	1.0000	1.0000	1.0000	1.0000	
rf	Random Forest Classifier	1.0000	1.0000	1.0000	1.0000	1.0000	1.0000	1.0000	
ada	Ada Boost Classifier	1.0000	1.0000	1.0000	1.0000	1.0000	1.0000	1.0000	
gbc	Gradient Boosting Classifier	1.0000	1.0000	1.0000	1.0000	1.0000	1.0000	1.0000	
xgboost	Extreme Gradient Boosting	1.0000	1.0000	1.0000	1.0000	1.0000	1.0000	1.0000	
lightgbm	Light Gradient Boosting Machine	1.0000	1.0000	1.0000	1.0000	1.0000	1.0000	1.0000	
lr	Logistic Regression	0.9995	1.0000	1.0000	0.9990	0.9995	0.9989	0.9989	
knn	K Neighbors Classifier	0.9984	1.0000	0.9980	0.9990	0.9985	0.9968	0.9968	
et	Extra Trees Classifier	0.9925	0.9998	0.9950	0.9912	0.9930	0.9850	0.9852	
lda	Linear Discriminant Analysis	0.9237	0.9998	1.0000	0.8764	0.9339	0.8449	0.8557	
ridge	Ridge Classifier	0.9232	0.0000	1.0000	0.8757	0.9335	0.8438	0.8548	
nb	Naive Bayes	0.9168	0.9986	1.0000	0.8667	0.9283	0.8306	0.8433	
qda	Quadratic Discriminant Analysis	0.9115	0.9982	0.9990	0.8598	0.9239	0.8196	0.8334	
dummy	Dummy Classifier	0.5360	0.5000	1.0000	0.5360	0.6979	0.0000	0.0000	

Logistic Regression is much more accurate and at the top, yet not perfect as in some of the other data models, however the accuracy is not reality.

```
1 print(predictions)
News Sentiment Change prediction_label prediction_score
0 -0.6166 -182.270004 0 1.0
1 0.2823 -17.070000 0 1.0
2 0.7861 -216.770004 0 1.0
3 0.5859 5.470000 1 1.0
4 0.3612 222.680006 1 1.0
5 -0.9584 314.640015 1 1.0
6 -0.9235 6.120000 1 1.0
7 -0.6788 -571.849976 0 1.0
8 -0.9830 -73.680000 0 1.0
9 -0.4939 -621.969971 0 1.0
10 0.9290 -775.580000 0 1.0
11 -0.9681 -0.790000 0 1.0
12 -0.8494 100.110001 1 1.0
13 0.5150 114.699997 1 1.0
14 0.7341 418.899994 1 1.0

[ ] 1 predictions.describe()
News Sentiment Change prediction_label prediction_score
count 15.000000 15.000000 15.000000 15.0
mean -0.159280 -51.824009 0.466667 1.0
std 0.753938 290.649231 0.516398 0.0
min -0.983000 -621.969971 0.000000 1.0
25% -0.901450 -199.520004 0.000000 1.0
50% -0.493900 -0.790000 0.000000 1.0
75% 0.530900 107.404999 1.000000 1.0
max 0.999000 418.899994 1.000000 1.0

[ ] 1 # Prediction ### 0,0,0,1,1,1,1,0,0,0,0,1,1,1
2 # Actual ### 0,0,1,1,1,1,0,0,0,0,1,1,1,0
3 # Results ### 1,1,0,1,1,1,0,1,1,1,1,0,1,1,0
4
5 11/15
6
7
8
9
10
0.7333333333333333
```

The accuracy of the results were much better, which was more focused on how the previous daily change of the DJIA affected the possible mood and direction of the next day. Even though still a small test dataset, I am more confidence in the modeling. We not have 11/15 accuracy, a 73% accuracy rating.

1 best = compare_models()

	Model	Accuracy	AUC	Recall	Prec.	F1	Kappa	MCC	TT (Sec)
dt	Decision Tree Classifier	1.0000	1.0000	1.0000	1.0000	1.0000	1.0000	1.0000	
svm	SVM - Linear Kernel	1.0000	0.0000	1.0000	1.0000	1.0000	1.0000	1.0000	
rf	Random Forest Classifier	1.0000	1.0000	1.0000	1.0000	1.0000	1.0000	1.0000	
ada	Ada Boost Classifier	1.0000	1.0000	1.0000	1.0000	1.0000	1.0000	1.0000	
gbc	Gradient Boosting Classifier	1.0000	1.0000	1.0000	1.0000	1.0000	1.0000	1.0000	
xgboost	Extreme Gradient Boosting	1.0000	1.0000	1.0000	1.0000	1.0000	1.0000	1.0000	
lightgbm	Light Gradient Boosting Machine	1.0000	1.0000	1.0000	1.0000	1.0000	1.0000	1.0000	
lr	Logistic Regression	0.9995	1.0000	1.0000	0.9990	0.9995	0.9989	0.9989	
knn	K Neighbors Classifier	0.9984	1.0000	0.9980	0.9990	0.9985	0.9968	0.9968	
et	Extra Trees Classifier	0.9925	0.9998	0.9950	0.9912	0.9930	0.9850	0.9852	
lda	Linear Discriminant Analysis	0.9237	0.9998	1.0000	0.8764	0.9339	0.8449	0.8557	
ridge	Ridge Classifier	0.9232	0.0000	1.0000	0.8757	0.9335	0.8438	0.8548	
nb	Naive Bayes	0.9168	0.9986	1.0000	0.8667	0.9283	0.8306	0.8433	
qda	Quadratic Discriminant Analysis	0.9115	0.9982	0.9990	0.8598	0.9239	0.8196	0.8334	
dummy	Dummy Classifier	0.5360	0.5000	1.0000	0.5360	0.6979	0.0000	0.0000	

[9] 1 predictions = predict_model(best, data=test)

Even though our second attempt in Machine Learning resulted in a much more accurate result of at least 73% accuracy, the overall potential is much more promising. A more weighted, calculated analysis in data preparation, feature engineering, and data modeling would significantly increase the accuracy. Run it through Deep Learning, the accuracy could be in the high 90%. Would it be better than the twitter 87% accuracy, I am unsure at this time.

In conclusion, this does not promise an absolute accurate prediction of how well the premarket newsfeed will forecast the daily close value. However, in regard to Chaos Theory, the premarket newsfeed, might be a strong suspect for the Strange Attractor, as it is dependent on the previous day's emotions. Thank you.

References and Sources

1. Chaos Theory

<https://medium.datadriveninvestor.com/chaos-theory-and-market-predictions-9e03ec47d67e>

2. Federal Reserve Bank of San Francisco

<https://www.frbsf.org/economic-research/indicators-data/daily-news-sentiment-index/>

3. MarketWatch – DJIA Historical Open Close

<https://www.marketwatch.com/investing/index/djia/download-data?startDate=7/5/2016&endDate=3/30/2017>

4. Sentiment Analysis for Stock Price Prediction in Python

<https://towardsdatascience.com/sentiment-analysis-for-stock-price-prediction-in-python-bed40c65d178>

5. A Short Introduction to VADER

<https://towardsdatascience.com/an-short-introduction-to-vader-3f3860208d53>

6. Understanding Lexicon Based Sentiment Analysis

<https://www.knime.com/blog/lexicon-based-sentiment-analysis>

7. NLP In the Stock Market

<https://towardsdatascience.com/nlp-in-the-stock-market-8760d062eb92>

